PushT, one planning step from one start state: mean and variance over seeds

PS / MPPI / DIAL: pusht.py hyperparameters, 2000 samples × 5 iterations from zero knots (10,005 rollouts); KSOS ×5: 5 restarts × 1000 two-level Sobol samples (box 0.8^i around the last solution), ADMM-cuda (5,005 rollouts); + LSE: 100 perturbations per sample (σ 0.2 → 0.01, λ = 1), then MPPI 1000 × 5 (σ 0.2; 505,015 rollouts); + GNC: costs kernel-smoothed over the restart's own samples (bandwidth 1.0σ → 0.1σ, no extra rollouts), then MPPI 1000 × 5 (σ 0.2; 10,010 rollouts); seeds 0–4 (n = 5); std uses ddof = 1

Plan cost: mean ± std over seeds (dots = seeds)

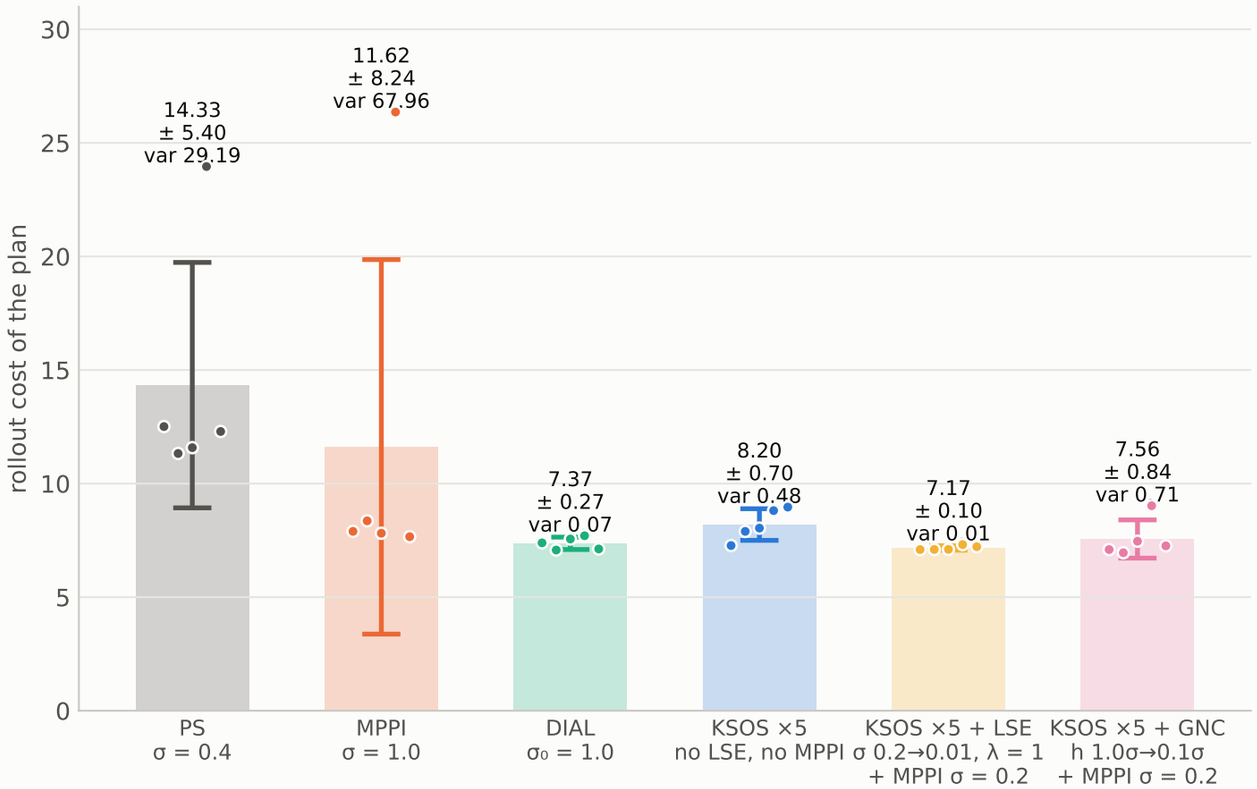
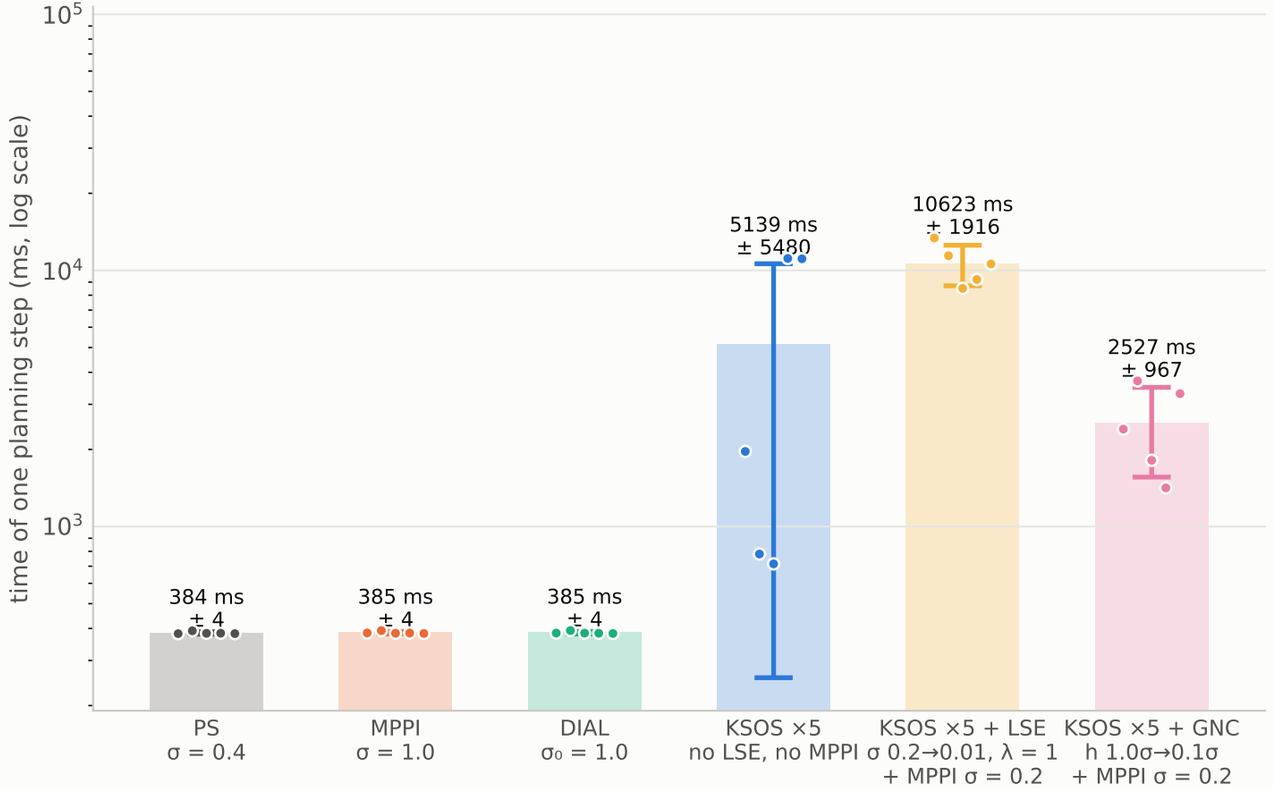
Planning time: mean ± std over seeds