PushT, one planning step from one start state (cold start): KSOS vs PS / MPPI / DIAL

PS / MPPI / DIAL: pusht.py hyperparameters with 1000 samples × 5 iterations; KSOS: 5 restarts × 1000 two-level Sobol samples (box 0.8^i around the last solution), ADMM-cuda, then MPPI 1000 × 5 (noise σ = 0.1 or 0.2, temp 0.1); seeds 0–4

Plan cost (dots = seeds, bar = mean; lower is better)

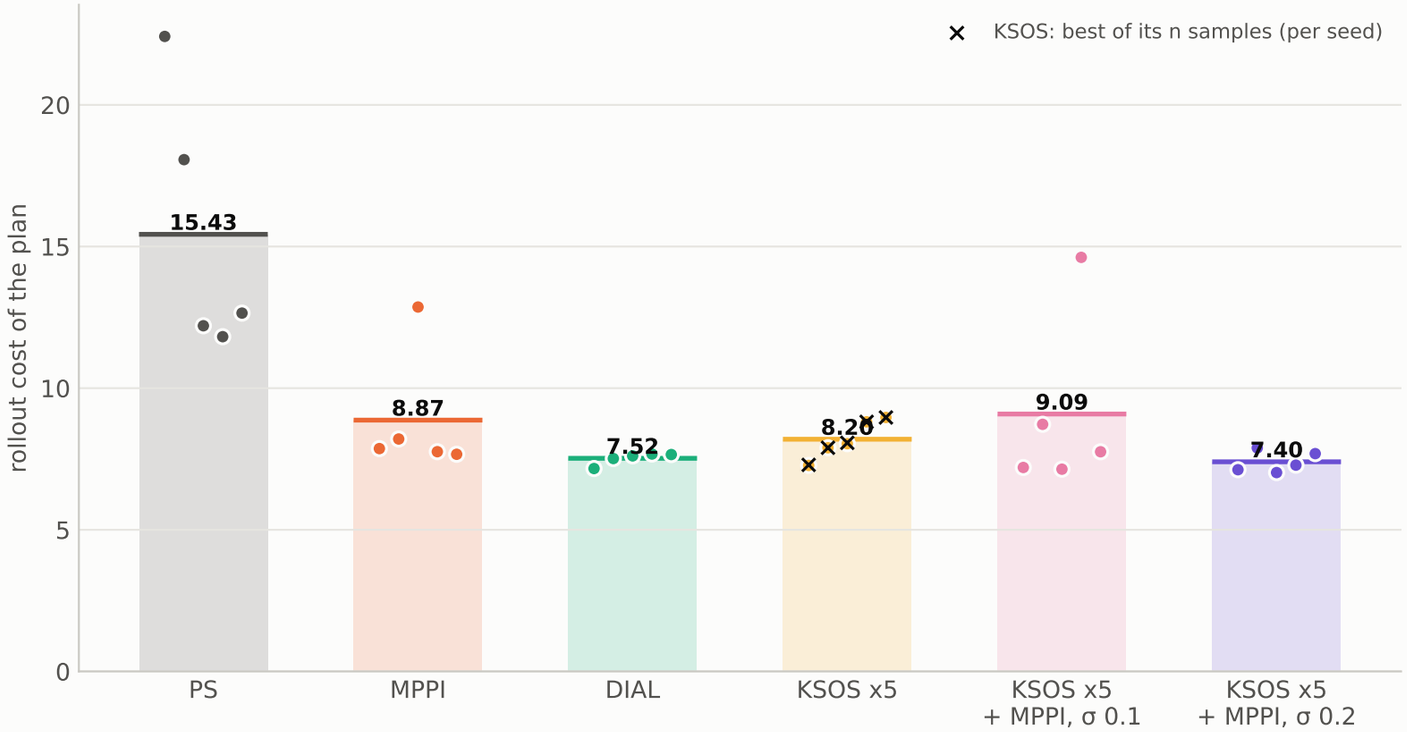
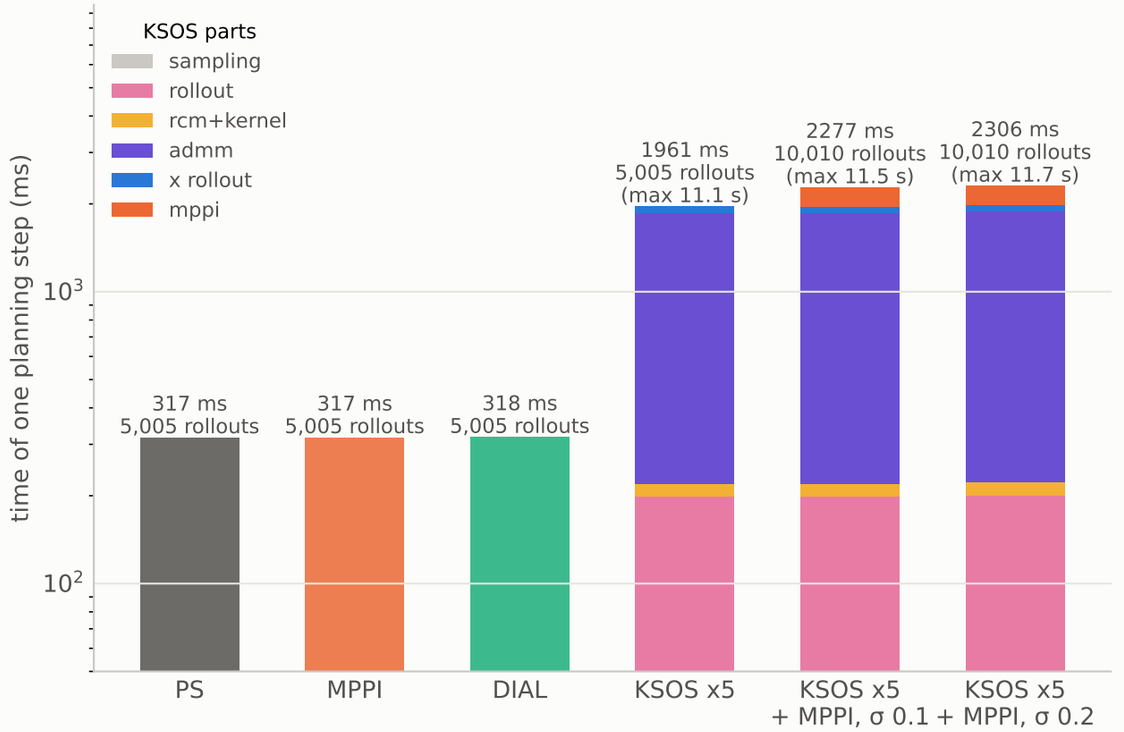
Planning time (median over seeds, log scale, JIT excluded)