PushT (12-D: 6 knots × 2 controls), two-level Sobol samples: n/10 clusters × 10 points, RCM order

one start state, seed 0; Wendland φ_{12} (newton-rs: Laplacian, same σ), λ = 1e-5; one timed run per point, each in a fresh process; ADMM eps 1e-6 / 1e-5; a solver that fails at one n is skipped for larger n

Solve time, converged runs only

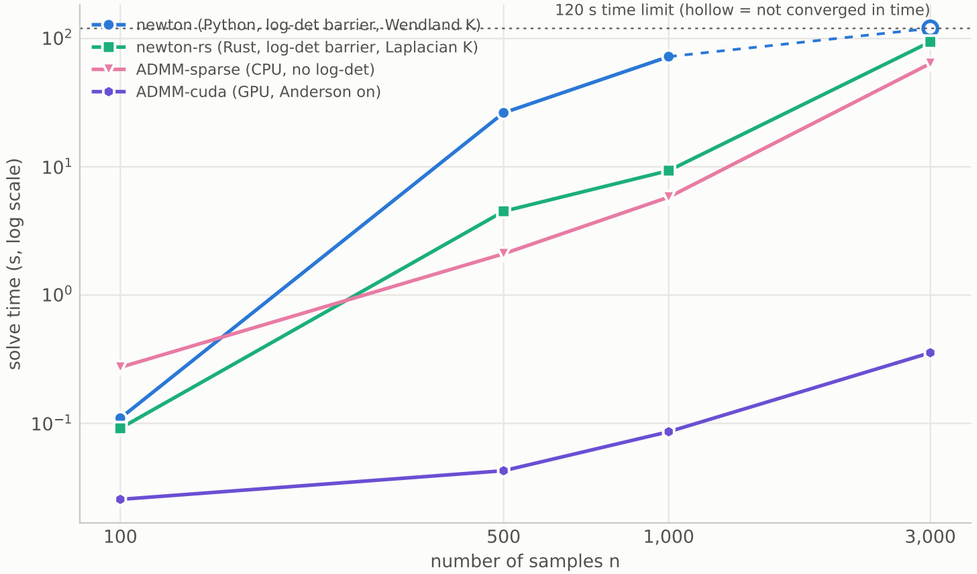

Bound c vs. best sample

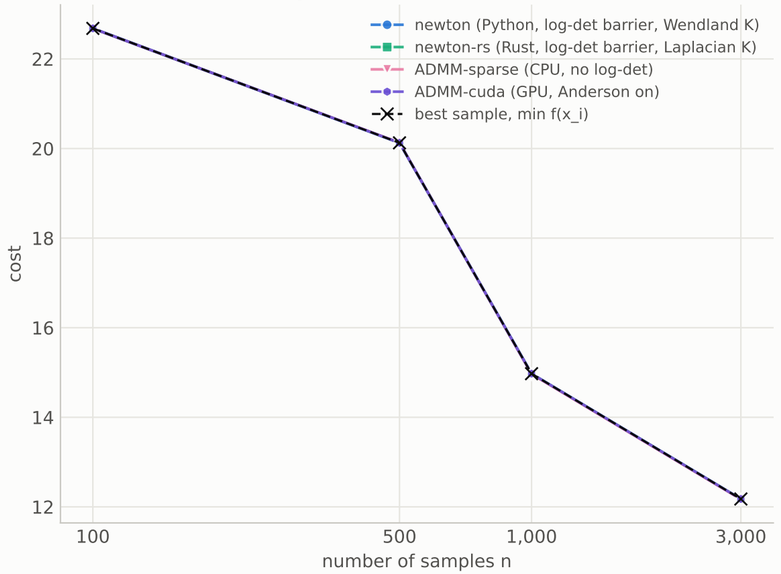
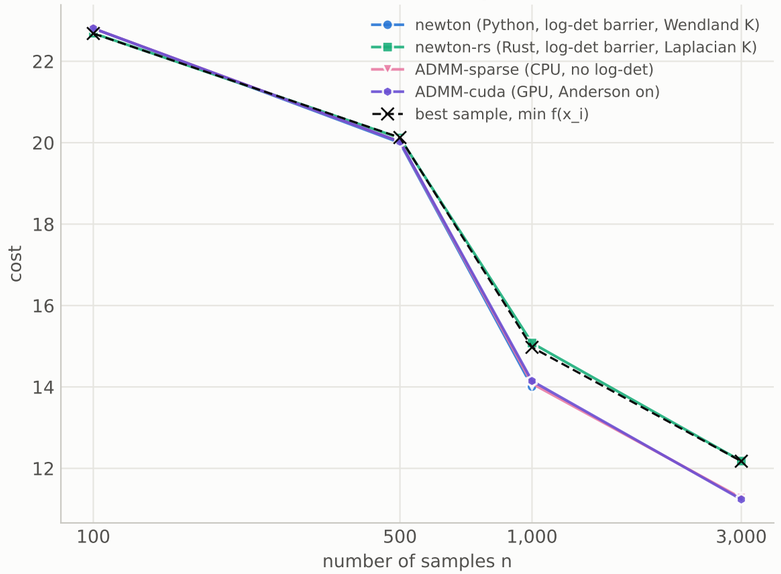
Rollout cost at x = Σ α_i x_i vs. best sample